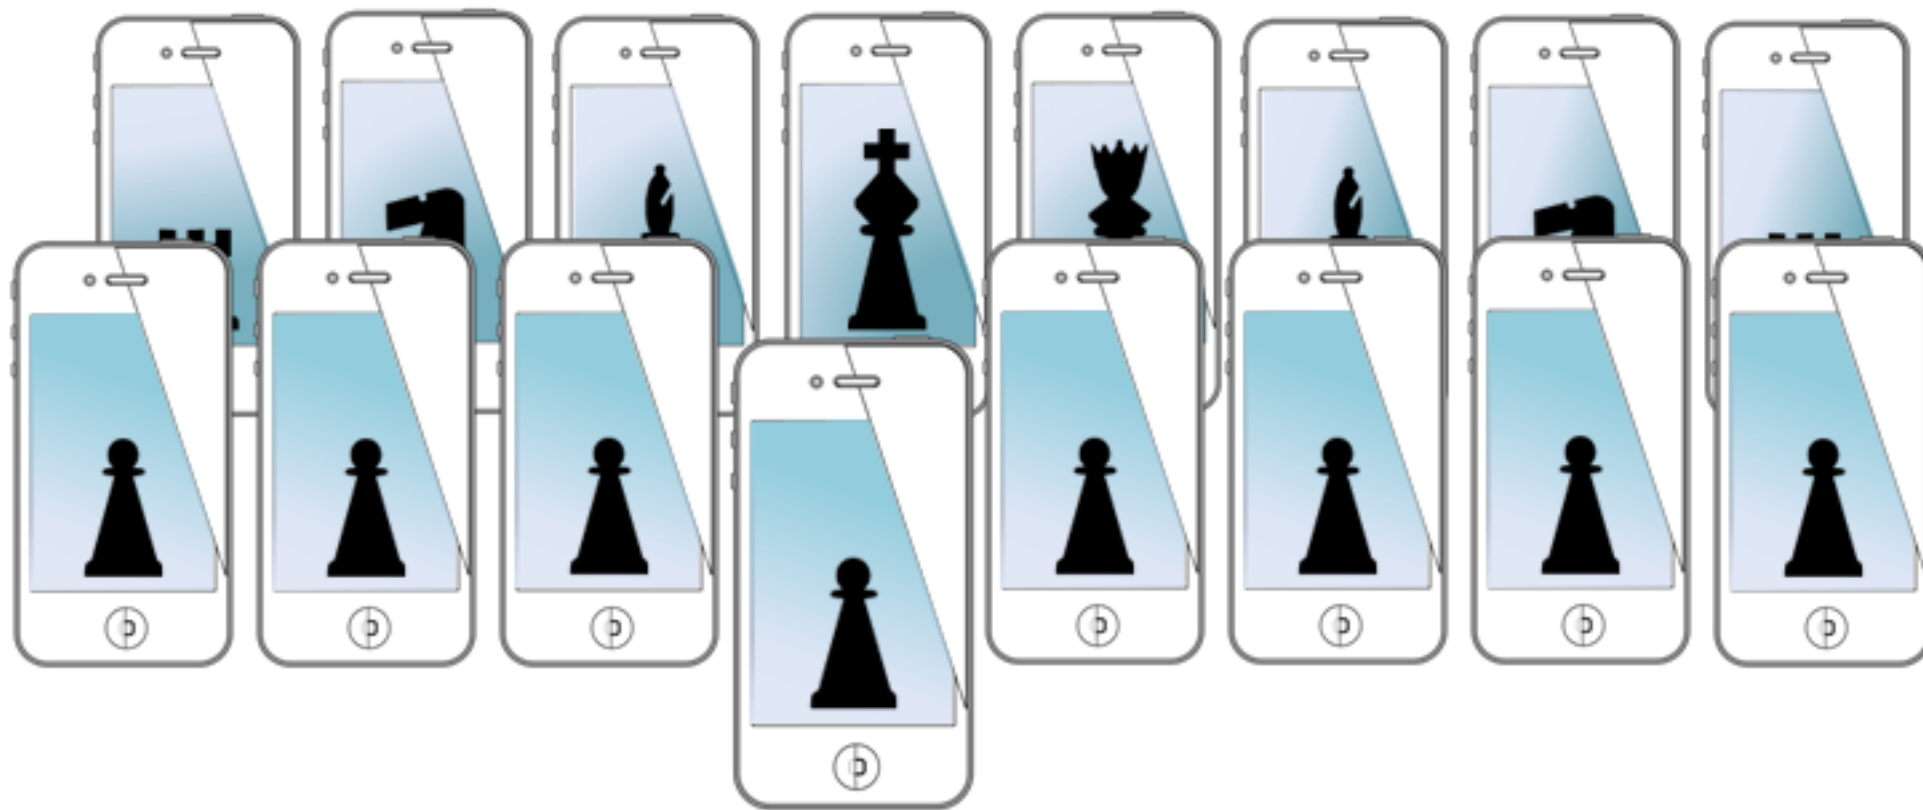


MOBILE SENSING LEARNING & CONTROL



CSE5323 & 7323

Mobile Sensing, Learning, and Control

lecture six: audio graphing, sampled data, accelerate, & FFT

Eric C. Larson, Lyle School of Engineering,
Computer Science and Engineering, Southern Methodist University

course logistics

- A1 grades by Wednesday (I hope)
- A2 is up!

Module A

Create an iOS application using the NovocaineExample template that:

- Reads from the microphone
- Takes an FFT of the incoming audio stream
- Displays the frequency of the two loudest tones within 6Hz accuracy
- Is able to distinguish tones as least 25Hz apart, lasting for 100ms or more

The sound source must be external to the phone (i.e., laptop, instrument, another phone, etc.).

Module B

Create an iOS application using the NovocaineExample template that:

- Reads from the microphone
- Plays a settable (via a slider or setter control) inaudible tone to the speakers (15-20kHz)
- Displays the magnitude of the FFT of the microphone data in decibels
- Is able to distinguish when the user is {not gesturing, gestures toward, or gesturing away} from the microphone using Doppler shifts in the frequency

agenda


- graphing audio fast (well, graphing anything)
 - must use lowest level graphing, OpenGL
- dealing with sampled data
- the accelerate framework
 - massive digital signal processing library
- the Fourier Transform for spectral analysis

audio graphing

- we want to see the incoming samples
 - good for debugging
 - equalizers
 - oscilloscope type applications

how much data to show?

- sampling at 44.1kHz == 44100 samples per second



0.5 seconds is
22050 samples

display is 640
pixels wide


what if we want
lots of graphs?

solution

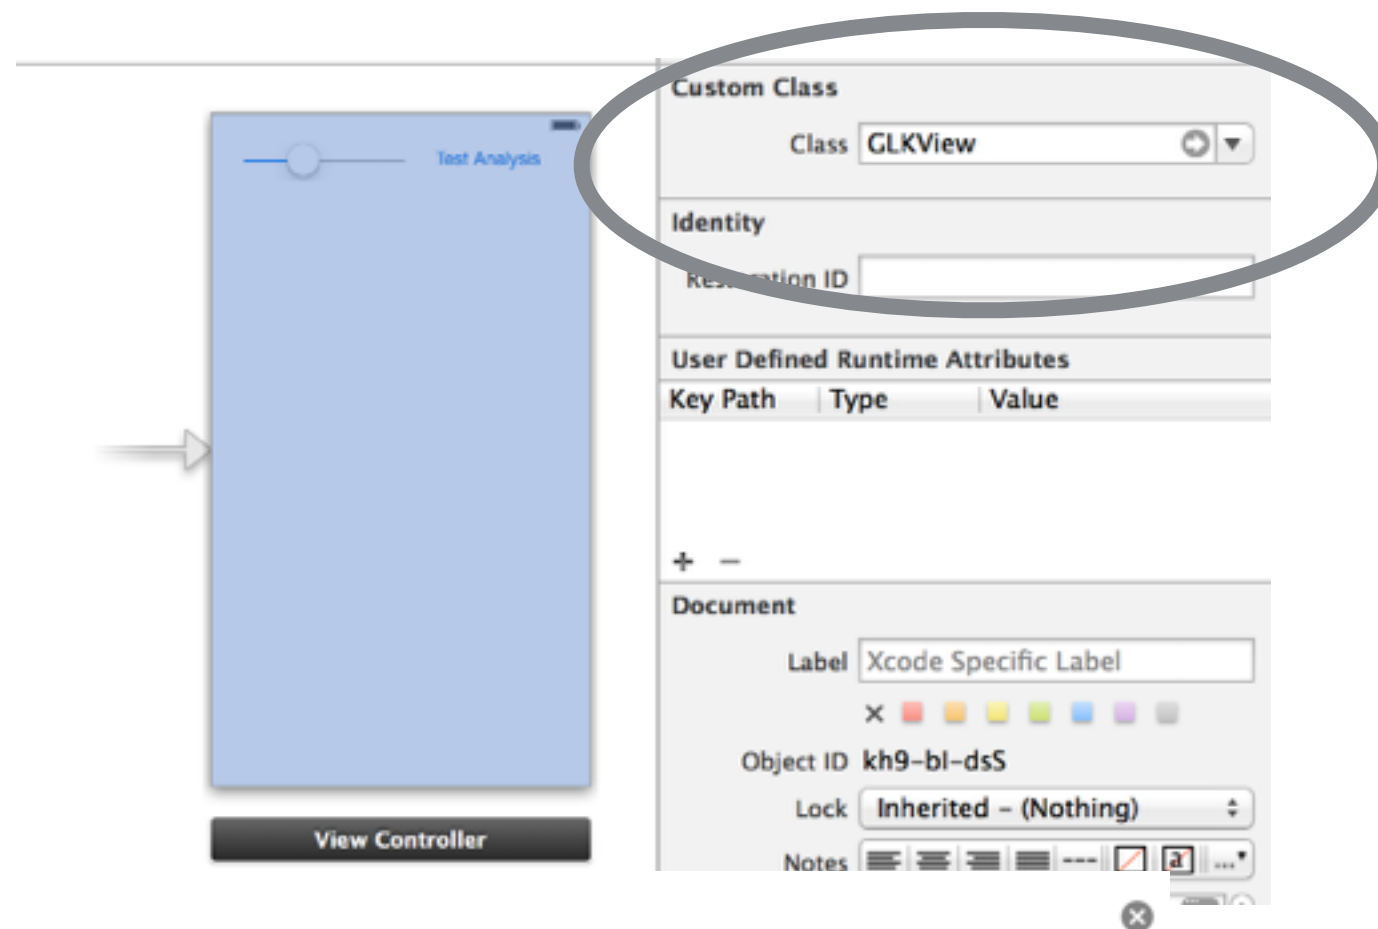
- use the GPU
- set vectors of data on a 2D plane
- let the renderer perform scaling, anti-aliasing, and bit blitting to screen
- ...this is not a graphics course

easy solution

- use graph helper, which uses GLKView and GLKViewController



add this to project



▼ Link Binary With Libraries (4 items)

Name	Status
 GLKit.framework	Required ▾

add GLKit framework

the graph helper

```
#import <GLKit/GLKit.h>
```

```
@interface YourCustomViewController : GLKViewController
```

```
GraphHelper *graphHelper;
```

declare in implementation

inherit from open GL

In View Did Load

```
// start animating the graph
```

```
int framesPerSecond = 15;
```

```
int numDataArraysToGraph = 3;
```

```
graphHelper = new GraphHelper(self,
```

```
framesPerSecond,
```

```
numDataArraysToGraph,
```

```
PlotStyleSeparated); //drawing starts immediately after call
```

setup GLKViewController

```
graphHelper->SetBounds(-0.9,0.9,-0.9,0.9); // bottom, top, left, right,
```

```
//full screen==(-1,1,-1,1)
```

bounds for screen

```
-(void)dealloc{
```

```
graphHelper->tearDownGL();
```

```
// ARC handles everything else, just clean up what we used c++ for (calloc, malloc, new)
```

```
}
```

```
enum PlotStyle {  
    PlotStyleOverlaid,  
    PlotStyleSeparated  
};
```


setting data

```
// override the GLKViewController draw function, from OpenGL ES
- (void)glkView:(GLKView *)view drawInRect:(CGRect)rect {
    graphHelper->draw(); // draw the graph
}
```

called for each draw to screen

```
void setGraphData(int graphNum, float *data, int dataLength, float normalization = 1.0,
                  float minValue = 0.0)
```

prototype for setting scatter data

```
// override the GLKView update function, from OpenGL ES
- (void)update{
```

set data for 0th graph

```
    graphHelper->setGraphData(0, //channel index
                              inputAudioDataBuffer, //data
                              kBufferSize/2.0, //data length
                              64.0); // max value to normalize (==1 if not set)
```

```
    // just plot the audio stream
```

```
    graphHelper->setGraphData(1, inputAudioDataBuffer, kBufferSize); // set graph channel
```

```
    graphHelper->update(); // update the graph
```

set data for 1st graph

update render state




audio graphing demo!

intro to sampled data

- why is understanding sampled data important?
 - because we'll be dealing with it all semester
 - it's important to understand basic mistakes that can be made
- there are entire courses dedicated to sampled time series
 - actually entire courses on analyzing frequency content
- we'll touch on a few guidelines to help you design your projects better

intro to sampled data


- physical processes are continuous
 - to process with computers, we must digitize it
 - digitization can change how we understand the signal
- digitization occurs in time and amplitude
 - time: sampling
 - amplitude: quantization



sampled data


- quantization (amplitude)
 - introduces error in estimating amplitude of a signal
 - error can be reduced by adding more “bits per sample”
- most ADCs are 16 bits, considered “good enough”
- sufficient for most uses
 - not for others!

iPhone uses LPCM 32 bits, Q8.24



sampling errors

- sampling in time
 - introduces errors through 'aliasing'
 - limits the range of frequencies able to be accurately captured
 - root of most common mistakes with sampled data



so how do I sample?

- heuristics
 - don't try to sample extremely small increments or values!
 - if capturing an "X"Hz signal, need to sample at least 2"X" Hz
 - changing sample rates is complicated, don't just drop every other sample
- for example, speech
 - majority of necessary energy in speech is located $< 8000\text{Hz}$
 - phones (for speech) typically capture at 16KHz or lower
 - good enough for speech, not music!

sanity check

- I need to detect an 80Hz signal
 - what sampling rate should we use?
- I want to detect a feather dropping next to the microphone
 - can the sound be detected?

the accelerate framework

- very powerful digital signal processing (DSP) library
 - look at vDSP Programming Guide on developer.apple.com for the complete API
- provides mathematics for performing fast DSP

input data stride scalar output array length

vDSP_vsmul(data, 1, &mult, data, 1, numFrames*numChannels);

```
void vDSP_vsmul (
    const float __vDSP_input1[],
    vDSP_Stride __vDSP_stride1,
    const float *__vDSP_input2,
    float __vDSP_result[],
    vDSP_Stride __vDSP_strideResult,
    vDSP_Length __vDSP_size
);
```

an example

```
[audioManager setInputBlock:^(float *data, UInt32 numFrames, UInt32 numChannels) {  
    float volume = userSetVolumeFromSlider;  
    vDSP_vsmul(data, 1, &volume, data, 1, numFrames*numChannels);  
    ringBuffer->AddNewInterleavedFloatData(data, numFrames, numChannels);  
}];
```

```
[audioManager setInputBlock:^(float *data, UInt32 numFrames, UInt32 numChannels) {  
  
    // get the max  
    float maxVal = 0.0;  
    vDSP_maxv(data, 1, &maxVal, numFrames*numChannels);  
  
    printf("Max Audio Value: %f\n", maxVal);  
  
}];
```

```
[audioManager setInputBlock:^(float *data, UInt32 numFrames, UInt32 numChannels)  
{  
    vDSP_vsq(data, 1, data, 1, numFrames*numChannels);  
    float meanVal = 0.0;  
    vDSP_meanv(data, 1, &meanVal, numFrames*numChannels);  
}];
```

processing audio


- lots of space to explore
 - great reference: “DSP First” by McClellan, Schafer, Yoder
 - <http://www.rose-hulman.edu/DSPFirst/visible3/contents/index.htm>
 - filtering
 - manipulate signal: high, low, bandpass
 - analysis
 - analyze characteristics of signal (like speech recognition)
 - synthesis
 - play around with different ideas, and see what sounds good!
 - not just pure synthesis, but also manipulation (like a guitar effect)
- for now, we’re going to stick with analysis
 - specifically, the **Fourier Transform**

the Fourier transform

- extremely useful, not just for signal junkies but also:
 - computer scientists, engineers, physicists, mathematicians, astronomers, oceanographers, health care professionals, etc.
- the Fourier Transform (FT) converts a time series into a frequency spectrum
 - the spectrum is an array of complex numbers which we will represent in polar form (i.e., with magnitude and phase)
 - each complex number represents a sinusoidal wave at a specific frequency
- we will use the FFT in the accelerate framework
 - complexity is $O(N \log_2(N))$ (for radix 2 FFT)

what is the FFT?

think of it as a vector projection (intuitively)



(4,2)

where did these numbers
come from?

$$\vec{x} \quad (1,0) \bullet (4,2) = 4$$

$$\vec{y} \quad (0,1) \bullet (4,2) = 2$$

point by point multiplication and add it up!

$$\vec{x} \bullet \vec{y} = 0$$

$$|\mathbf{x}| = |\mathbf{y}| = 1$$


tells us “how much” of the vector

is the results of other orthogonal vectors

\vec{x}
 \vec{y}

what is the FFT?



what if the orthogonal vectors were functions?



the FFT


frequency content	0f	0.7	0°				= 0.7
	1f	2.5	10°	+	2.5	-10°	-1f ~ 2.5 cos(2pi (f) t+10°)
	2f	0.1	45°	+	0.1	-45°	-2f ~ 0.1 cos(2pi (2f) t+45°)
	3f	0.5	-12°	+	0.5	12°	-3f ~ 0.5 cos(2pi (3f) t-12°)
		⋮	⋮		⋮	⋮	
	N f	.005	0°	+	.005	0°	-N f ~ 0.005 cos(2pi (Nf) t)

time and frequency

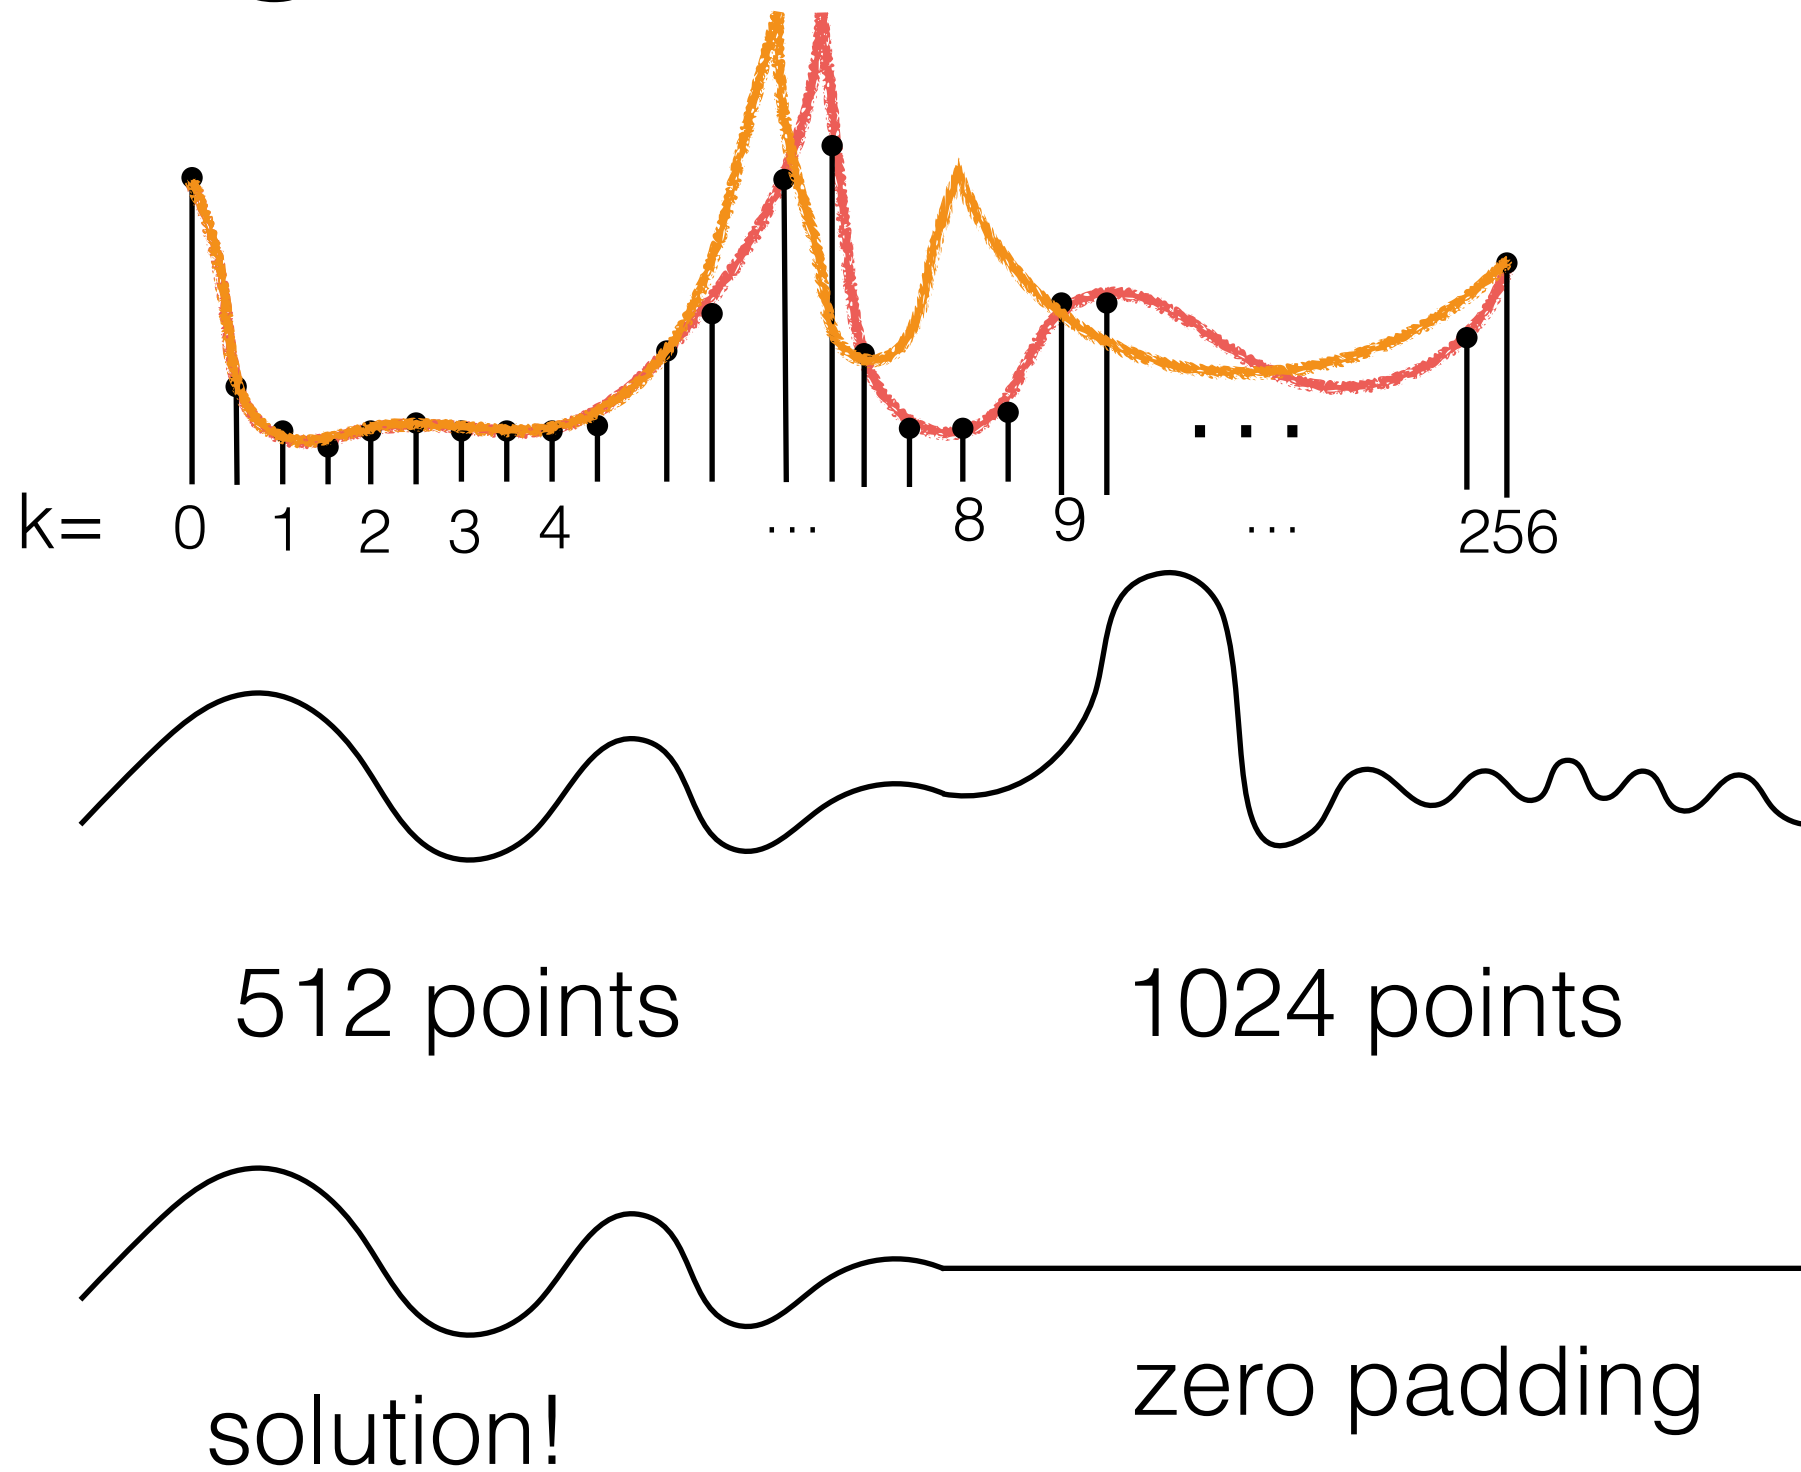


magnitude spectrum

time and frequency




using the FFT



using the FFT

raw audio




magnitude FFT



magnitude FFT
in dB

$20 \log_{10}(|\text{FFT}|)$





100

0

40

-30



programming the FFT

```
#import "SMUFFTHelper.h"
```

```
float          *fftMagnitudeBuffer;  
float          *fftPhaseBuffer;  
SMUFFTHelper  *fftHelper;
```

fft size

window size

window type

```
//setup the fft
```

```
fftHelper = new SMUFFTHelper(kBufferLength, kBufferLength, WindowTypeRect);
```

```
fftMagnitudeBuffer = (float *)calloc(kBufferLength/2, sizeof(float));
```

```
fftPhaseBuffer      = (float *)calloc(kBufferLength/2, sizeof(float));
```

```
free(fftMagnitudeBuffer);
```

```
free(fftPhaseBuffer);
```

```
delete fftHelper;
```

tear down in dealloc

```
enum WindowType {  
    WindowTypeHann,  
    WindowTypeHamming,  
    WindowTypeRect,  
    WindowTypeBlackman,  
};
```

```
fftHelper->forward(0, inputAudioDataBuffer, fftMagnitudeBuffer, fftPhaseBuffer);
```

reserved

input array

magnitude out


phase out

highly optimized!! even using tricks we have not discussed

for next time...

- programming the FFT
- filtering
- more frequency analysis with the FFT
 - windowing effects

MOBILE SENSING LEARNING & CONTROL




CSE5323 & 7323

Mobile Sensing, Learning, and Control

lecture six: audio graphing, sampled data, accelerate, & FFT

Eric C. Larson, Lyle School of Engineering,
Computer Science and Engineering, Southern Methodist University

MOBILE SENSING LEARNING & CONTROL



CSE5323 & 7323

Mobile Sensing, Learning, and Control

lecture seven: filtering and windowing

Eric C. Larson, Lyle School of Engineering,
Computer Science and Engineering, Southern Methodist University

course logistics


- A1 grades posted
 - feedback from blackboard system
 - team member who turned in A1 has the feedback.
- A2 is due Friday of next week
 - constraints updated for 12Hz accuracy ($\pm 6\text{Hz}$)

agenda

- using the FFT
- we'll be taking a brief look at filtering
 - an in-depth treatment would take many weeks, we will just do the basics
- how to use a filter (FIR)
 - how to create the most basic of filters
 - two methods of applying the filter to a signal
- We will NOT cover filter design
 - For the curious, the only filter design we'll touch is windowed FIR


some fft examples

1kHz sine wave





some fft examples

1kHz sine wave + 2.7kHz sine wave



some fft examples



a person whistling



frequency (Hz)

some fft examples



humming



frequency (Hz)

some fft examples

humming again



frequency (Hz)

fft demo!!

```
#import "SMUFFTHelper.h"
```

```
float          *fftMagnitudeBuffer;  
float          *fftPhaseBuffer;  
SMUFFTHelper   *fftHelper;
```

fft size

window size

window type

```
//setup the fft
```

```
fftHelper = new SMUFFTHelper(kBufferLength, kBufferLength, WindowTypeRect);
```

```
fftMagnitudeBuffer = (float *)calloc(kBufferLength/2, sizeof(float));
```

```
fftPhaseBuffer      = (float *)calloc(kBufferLength/2, sizeof(float));
```

```
free(fftMagnitudeBuffer);
```

```
free(fftPhaseBuffer);
```

```
delete fftHelper;
```

tear down in dealloc

```
enum WindowType {  
    WindowTypeHann,  
    WindowTypeHamming,  
    WindowTypeRect,  
    WindowTypeBlackman,  
};
```

```
fftHelper->forward(0, inputAudioDataBuffer, fftMagnitudeBuffer, fftPhaseBuffer);
```

reserved


input array

magnitude out


phase out

noise in the FFT


- variance around actual magnitude unavoidable




local peak finding




peak interpolation




peak interpolation





peak interpolation




the doppler effect



the doppler effect



the doppler effect




the doppler effect

The diagram shows the equation for the Doppler effect, $\Delta f = \frac{V_{object}}{c} f_0$, with four callout boxes pointing to its components:

- A box labeled "change in frequency" points to Δf .
- A box labeled "velocity of object" points to V_{object} .
- A box labeled "speed of sound" points to c .
- A box labeled "frequency of source" points to f_0 .

$$\Delta f = \frac{V_{object}}{c} f_0$$



linear

db

db zoomed

filters!

- we will cover what we can...

signals and systems

- signals are collections of sampled data (arrays)
 - such as audio, accelerometer, etc.
 - can also be 2D, like images
- systems are objects which manipulate signals
 - characterized by their “input/output” relationships
 - we say “ $x[n]$ is passed through H , resulting in $y[n]$ ”




filters

- filters are systems which manipulate frequencies
 - certain frequencies to pass through, but not others
 - “lowpass” filter allows low frequencies to pass through
 - “highpass” filter likewise allows high frequencies through
- keep in mind: no filter is perfect!!
 - no filter will pass everything you want while stopping everything you don't
 - everything is a balance between different parameters you can control
- we won't study how to design filters
 - we will study properties of filters and how to use them
 - so we need to know what filters can and cannot do


filters in frequency

- filters can be characterized a few different ways
 - let's start by looking at their properties in the frequency domain
- filters have the following frequency-domain attributes:
 - passband gain
 - passband bandwidth
 - stopband attenuation
 - transition bandwidth



filters in time

- filter are also signals (time series)
 - the series is called the “impulse response” of the filter
 - the frequency-domain plots are just Fourier transforms of the impulse response (magnitude)
- the time-domain property we care about is length
 - everything else is best left to a filter design course



so how to design a filter?

- scipy.signal in python (try to use remez)
- decent tutorial:
 - <http://mpastell.com/2010/01/18/fir-with-scipy/>
- matlab
 - fdatool
- lots of other places!


```
from pylab import *  
import scipy.signal as signal  
n = 61  
a = signal.firwin(n, cutoff = 0.3,  
                  window = "hamming")
```


filtering by convolution



- we apply a filter using **convolution**
 - convolution allows us to combine frequency properties of two signals without taking an FFT
- basic principle:
 - convolution in time is multiplication in frequency
 - so the filter's frequency response will be multiplied by the frequency response of the signal

$$y[n] = \sum_{i=0}^{N-1} h[n-i]x[i]$$

convolution



$$y[n] = \sum_{i=0}^{N-1} h[n-i]x[i]$$



length


N

M

$N+M-1$

convolution efficiency


- algorithmic complexity
 - convolution is not particularly efficient, $O(N \times M)$
- to convolve faster, use that Fourier property:
 - “convolution in time is multiplication in frequency”
- why not just multiply to begin with!



its circular



- just using N point FFT performs “Circular Convolution”
 - which is not linear convolution
 - causes the tail end of the convolution to “wrap around” to the beginning
 - FFT assumes the function is periodic (we did not talk about this)
- be aware of circularity when filtering your signal with the FFT
 - zero-padding can solve this for you!
 - zero-pad both signals to a length that will contain the entire convolution, $N+M-1$
 - for streaming, you must use overlap-and-add!
 - http://en.wikipedia.org/wiki/Overlap-add_method

overlap and add







windowing: spectral BW widening





- multiplication in time is convolution in frequency
- a window is something we multiply in time with our signal
- windowing is unavoidable
 - why? we cannot take an infinite FFT...



windowing: spectral leakage



windowing: spectral leakage




which window to use?

- depends
 - narrowest main lobe: rect
 - good tradeoff: hamming (or Von Hann)
 - optimal tradeoff for a given bandwidth:
 - discrete prolate spheroidal sequence (dpss, Slepian taper)

for next time...

- core motion
 - the M7 co-processor

MOBILE SENSING LEARNING & CONTROL



CSE5323 & 7323

Mobile Sensing, Learning, and Control

lecture seven: filtering and windowing

Eric C. Larson, Lyle School of Engineering,
Computer Science and Engineering, Southern Methodist University