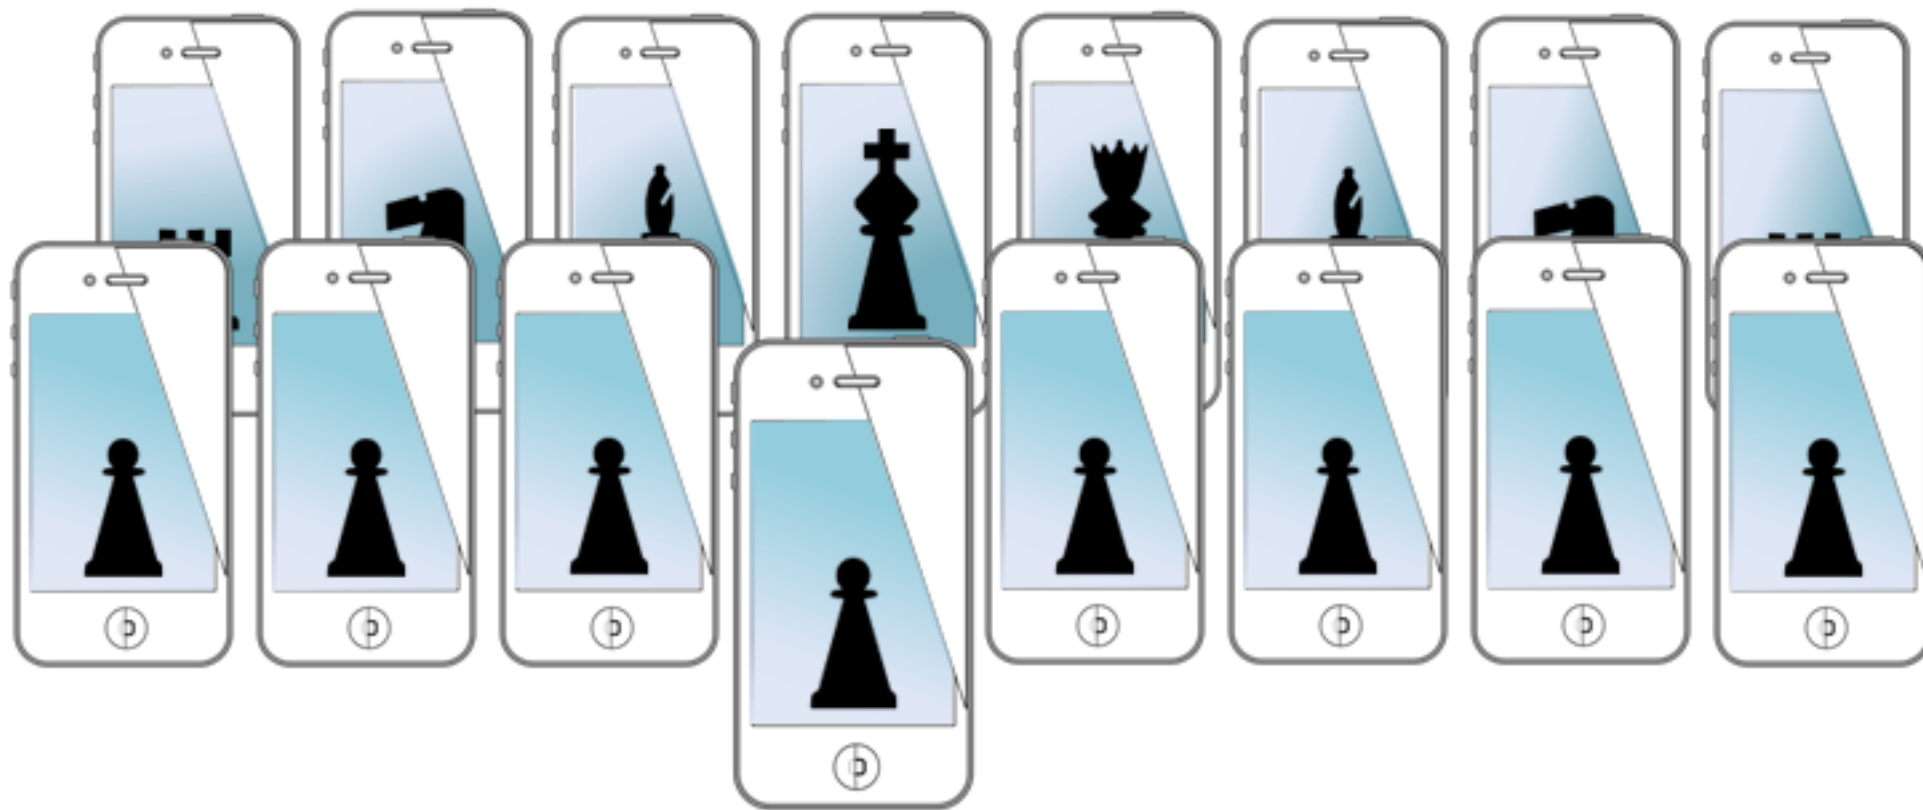


MOBILE SENSING LEARNING & CONTROL



CSE5323 & 7323

Mobile Sensing, Learning, and Control

lecture six: audio graphing, sampled data, accelerate, & FFT

Eric C. Larson, Lyle School of Engineering,
Computer Science and Engineering, Southern Methodist University

course logistics

- A1 grades by Wednesday (I hope)
- A2 is up!

Module A

Create an iOS application using the NovocaineExample template that:

- Reads from the microphone
- Takes an FFT of the incoming audio stream
- Displays the frequency of the two loudest tones within 6Hz accuracy
- Is able to distinguish tones as least 25Hz apart, lasting for 100ms or more

The sound source must be external to the phone (i.e., laptop, instrument, another phone, etc.).

Module B

Create an iOS application using the NovocaineExample template that:

- Reads from the microphone
- Plays a settable (via a slider or setter control) inaudible tone to the speakers (15-20kHz)
- Displays the magnitude of the FFT of the microphone data in decibels
- Is able to distinguish when the user is {not gesturing, gestures toward, or gesturing away} from the microphone using Doppler shifts in the frequency

agenda


- graphing audio fast (well, graphing anything)
 - must use lowest level graphing, OpenGL
- dealing with sampled data
- the accelerate framework
 - massive digital signal processing library
- the Fourier Transform for spectral analysis

audio graphing

- we want to see the incoming samples
 - good for debugging
 - equalizers
 - oscilloscope type applications

how much data to show?

- sampling at 44.1kHz == 44100 samples per second



0.5 seconds is
22050 samples

display is 640
pixels wide


what if we want
lots of graphs?

solution

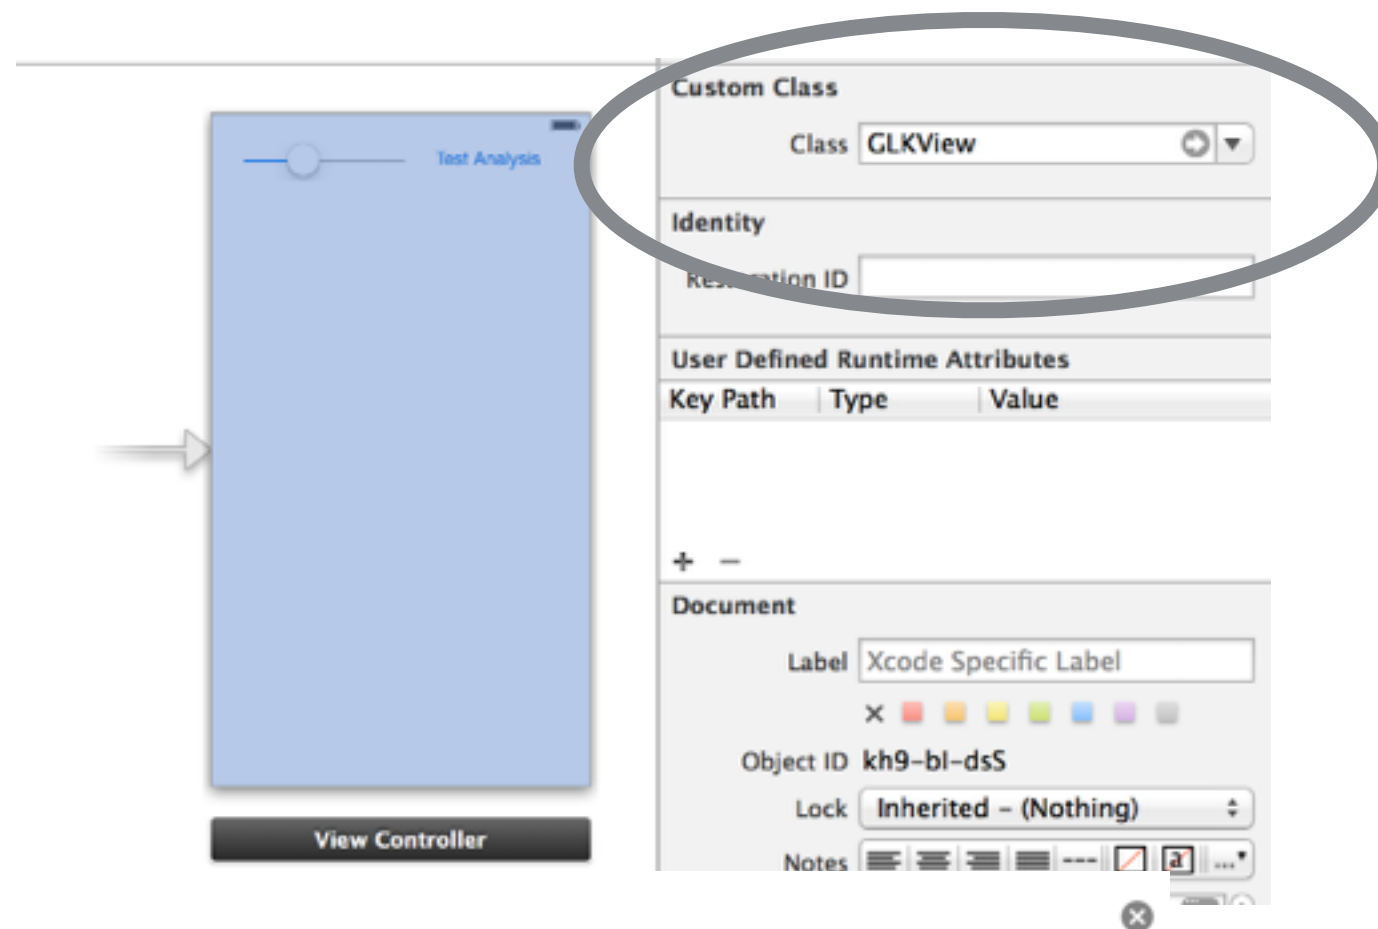
- use the GPU
- set vectors of data on a 2D plane
- let the renderer perform scaling, anti-aliasing, and bit blitting to screen
- ...this is not a graphics course

easy solution

- use graph helper, which uses GLKView and GLKViewController



add this to project



▼ Link Binary With Libraries (4 items)

Name	Status
 GLKit.framework	Required ▾

add GLKit framework

the graph helper

```
#import <GLKit/GLKit.h>
```

```
@interface YourCustomViewController : GLKViewController
```

```
GraphHelper *graphHelper;
```

declare in implementation

inherit from open GL

In View Did Load

```
// start animating the graph
```

```
int framesPerSecond = 15;
```

```
int numDataArraysToGraph = 3;
```

```
graphHelper = new GraphHelper(self,  
    framesPerSecond,  
    numDataArraysToGraph,  
    PlotStyleSeparated); //drawing starts immediately after call
```

setup GLKViewController

```
graphHelper->SetBounds(-0.9,0.9,-0.9,0.9); // bottom, top, left, right,  
//full screen==(-1,1,-1,1)
```

bounds for screen

```
-(void)dealloc{
```

```
graphHelper->tearDownGL();
```

```
// ARC handles everything else, just clean up what we used c++ for (calloc, malloc, new)
```

```
}
```

```
enum PlotStyle {  
    PlotStyleOverlaid,  
    PlotStyleSeparated  
};
```


setting data

```
// override the GLKViewController draw function, from OpenGL ES
- (void)glkView:(GLKView *)view drawInRect:(CGRect)rect {
    graphHelper->draw(); // draw the graph
}
```

called for each draw to screen

```
void setGraphData(int graphNum, float *data, int dataLength, float normalization = 1.0,
                  float minValue = 0.0)
```

prototype for setting scatter data

```
// override the GLKView update function, from OpenGL ES
- (void)update{
```

set data for 0th graph

```
    graphHelper->setGraphData(0, //channel index
                              inputAudioDataBuffer, //data
                              kBufferSize/2.0, //data length
                              64.0); // max value to normalize (==1 if not set)
```

```
    // just plot the audio stream
```

```
    graphHelper->setGraphData(1, inputAudioDataBuffer, kBufferSize); // set graph channel
```

```
    graphHelper->update(); // update the graph
```

set data for 1st graph

update render state




audio graphing demo!

intro to sampled data

- why is understanding sampled data important?
 - because we'll be dealing with it all semester
 - it's important to understand basic mistakes that can be made
- there are entire courses dedicated to sampled time series
 - actually entire courses on analyzing frequency content
- we'll touch on a few guidelines to help you design your projects better

intro to sampled data


- physical processes are continuous
 - to process with computers, we must digitize it
 - digitization can change how we understand the signal
- digitization occurs in time and amplitude
 - time: sampling
 - amplitude: quantization



sampled data


- quantization (amplitude)
 - introduces error in estimating amplitude of a signal
 - error can be reduced by adding more “bits per sample”
- most ADCs are 16 bits, considered “good enough”
- sufficient for most uses
 - not for others!

iPhone uses LPCM 32 bits, Q8.24



sampling errors

- sampling in time
 - introduces errors through 'aliasing'
 - limits the range of frequencies able to be accurately captured
 - root of most common mistakes with sampled data



so how do I sample?

- heuristics
 - don't try to sample extremely small increments or values!
 - if capturing an "X"Hz signal, need to sample at least 2"X" Hz
 - changing sample rates is complicated, don't just drop every other sample
- for example, speech
 - majority of necessary energy in speech is located $< 8000\text{Hz}$
 - phones (for speech) typically capture at 16KHz or lower
 - good enough for speech, not music!

sanity check

- I need to detect an 80Hz signal
 - what sampling rate should we use?
- I want to detect a feather dropping next to the microphone
 - can the sound be detected?

the accelerate framework

- very powerful digital signal processing (DSP) library
 - look at vDSP Programming Guide on developer.apple.com for the complete API
- provides mathematics for performing fast DSP

input data stride scalar output array length

vDSP_vsmul(data, 1, &mult, data, 1, numFrames*numChannels);

```
void vDSP_vsmul (
    const float __vDSP_input1[],
    vDSP_Stride __vDSP_stride1,
    const float *__vDSP_input2,
    float __vDSP_result[],
    vDSP_Stride __vDSP_strideResult,
    vDSP_Length __vDSP_size
);
```

an example

```
[audioManager setInputBlock:^(float *data, UInt32 numFrames, UInt32 numChannels) {  
    float volume = userSetVolumeFromSlider;  
    vDSP_vsmul(data, 1, &volume, data, 1, numFrames*numChannels);  
    ringBuffer->AddNewInterleavedFloatData(data, numFrames, numChannels);  
}];
```

```
[audioManager setInputBlock:^(float *data, UInt32 numFrames, UInt32 numChannels) {  
  
    // get the max  
    float maxVal = 0.0;  
    vDSP_maxv(data, 1, &maxVal, numFrames*numChannels);  
  
    printf("Max Audio Value: %f\n", maxVal);  
  
}];
```

```
[audioManager setInputBlock:^(float *data, UInt32 numFrames, UInt32 numChannels)  
{  
    vDSP_vsq(data, 1, data, 1, numFrames*numChannels);  
    float meanVal = 0.0;  
    vDSP_meanv(data, 1, &meanVal, numFrames*numChannels);  
}];
```

processing audio


- lots of space to explore
 - great reference: “DSP First” by McClellan, Schafer, Yoder
 - <http://www.rose-hulman.edu/DSPFirst/visible3/contents/index.htm>
 - filtering
 - manipulate signal: high, low, bandpass
 - analysis
 - analyze characteristics of signal (like speech recognition)
 - synthesis
 - play around with different ideas, and see what sounds good!
 - not just pure synthesis, but also manipulation (like a guitar effect)
- for now, we’re going to stick with analysis
 - specifically, the **Fourier Transform**

the Fourier transform

- extremely useful, not just for signal junkies but also:
 - computer scientists, engineers, physicists, mathematicians, astronomers, oceanographers, health care professionals, etc.
- the Fourier Transform (FT) converts a time series into a frequency spectrum
 - the spectrum is an array of complex numbers which we will represent in polar form (i.e., with magnitude and phase)
 - each complex number represents a sinusoidal wave at a specific frequency
- we will use the FFT in the accelerate framework
 - complexity is $O(N \log_2(N))$ (for radix 2 FFT)

what is the FFT?

think of it as a vector projection (intuitively)



(4,2)

where did these numbers
come from?

$$\vec{x} \quad (1,0) \bullet (4,2) = 4$$

$$\vec{y} \quad (0,1) \bullet (4,2) = 2$$

point by point multiplication and add it up!

$$\vec{x} \bullet \vec{y} = 0$$

$$|\mathbf{x}| = |\mathbf{y}| = 1$$


tells us “how much” of the vector

is the results of other orthogonal vectors

\vec{x}
 \vec{y}

what is the FFT?



what if the orthogonal vectors were functions?



the FFT


frequency content	0f	0.7	0°				= 0.7
	1f	2.5	10°	+	2.5	-10°	-1f ~ 2.5 cos(2pi (f) t+10°)
	2f	0.1	45°	+	0.1	-45°	-2f ~ 0.1 cos(2pi (2f) t+45°)
	3f	0.5	-12°	+	0.5	12°	-3f ~ 0.5 cos(2pi (3f) t-12°)
		⋮	⋮		⋮	⋮	
	N f	.005	0°	+	.005	0°	-N f ~ 0.005 cos(2pi (Nf) t)

time and frequency

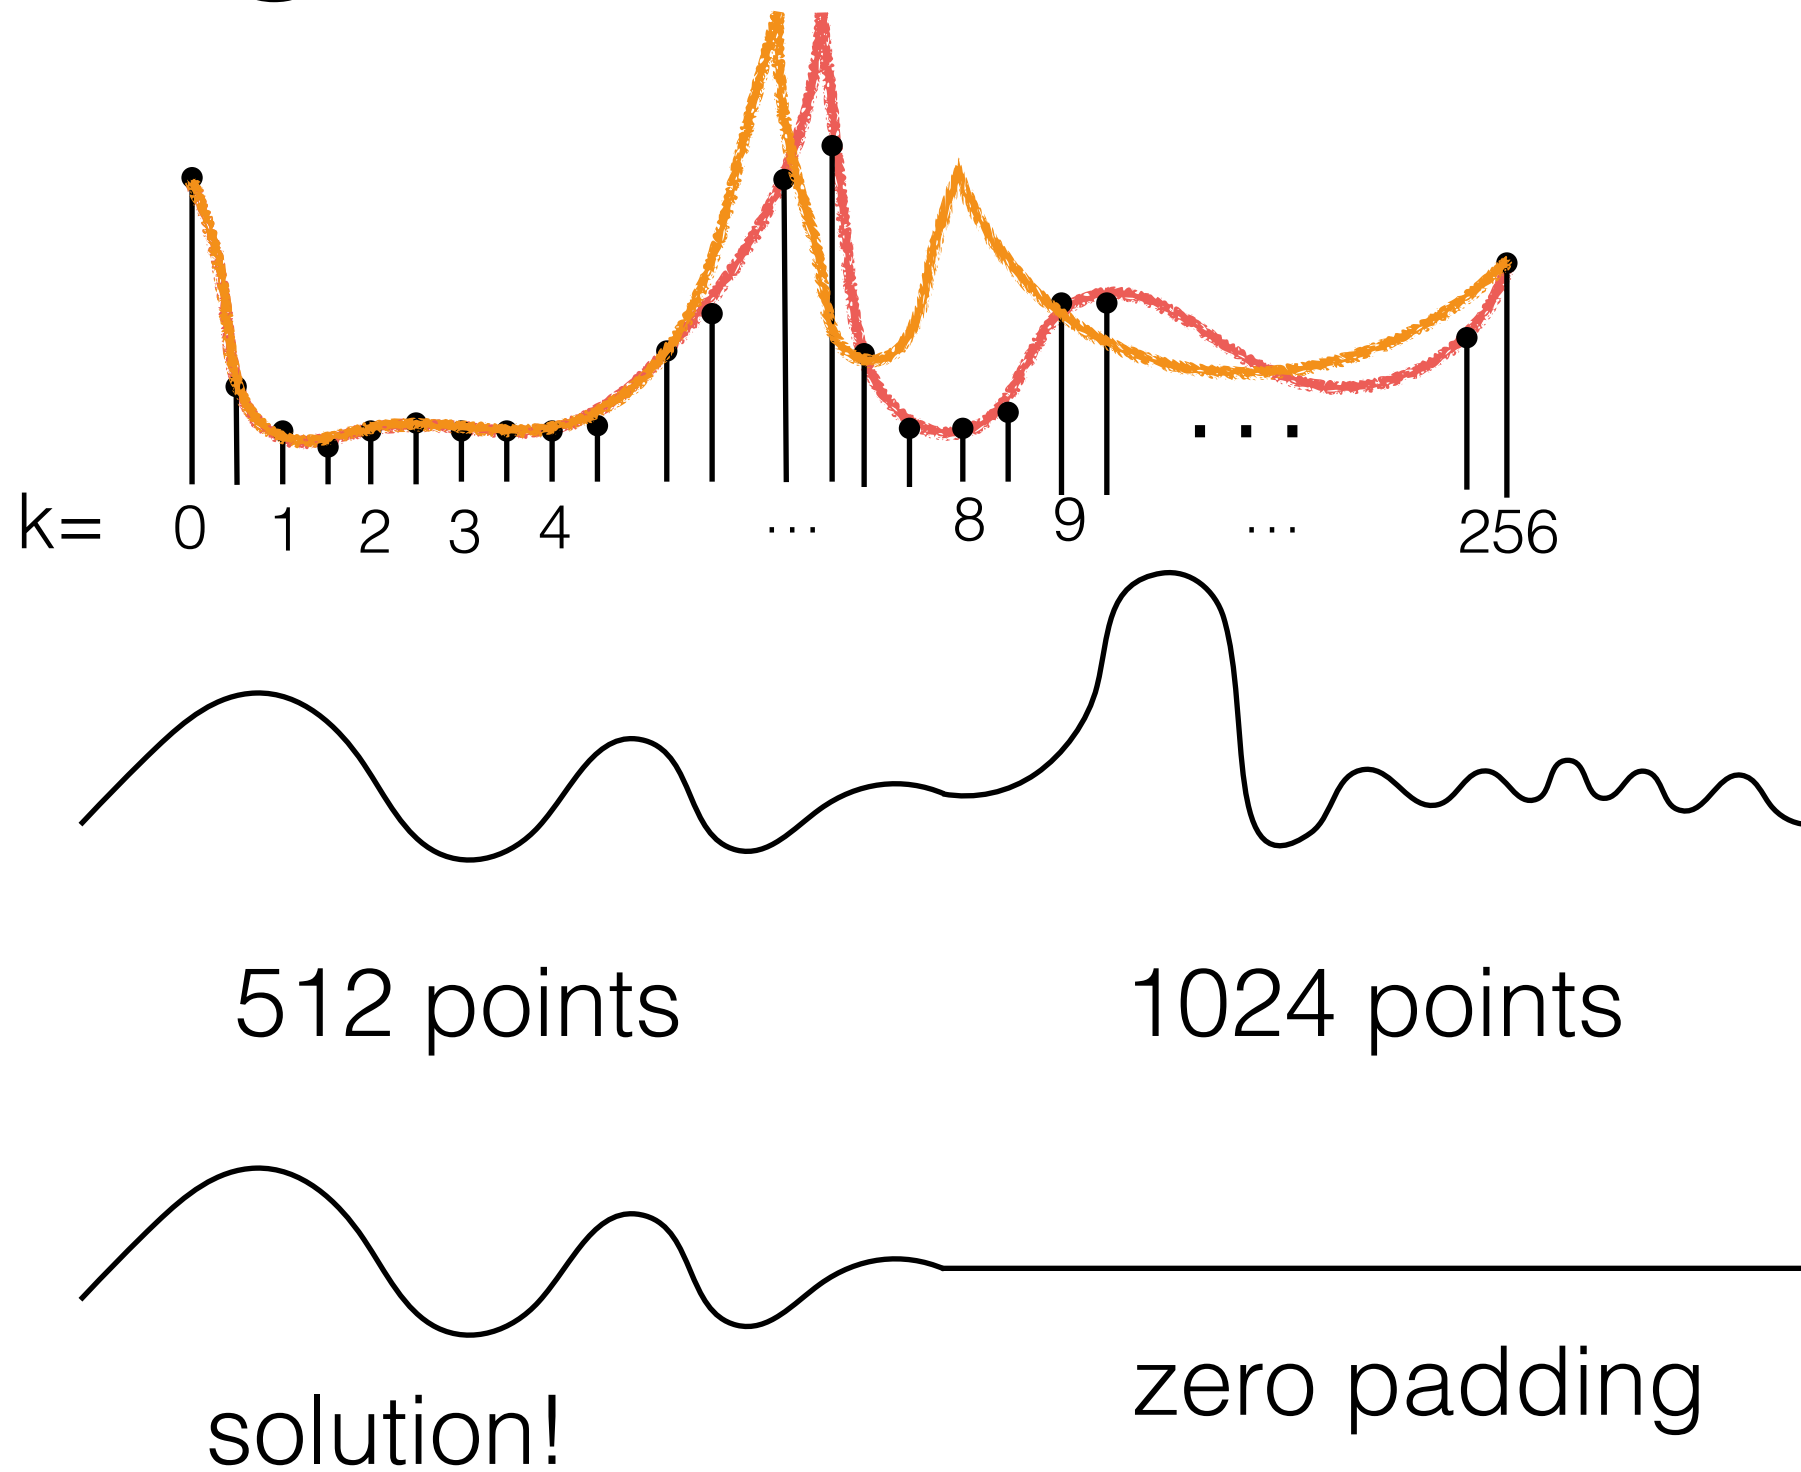


magnitude spectrum

time and frequency




using the FFT



using the FFT

raw audio




magnitude FFT



magnitude FFT
in dB

$20 \log_{10}(|\text{FFT}|)$





100

0

40

-30



programming the FFT

```
#import "SMUFFTHelper.h"
```

```
float          *fftMagnitudeBuffer;  
float          *fftPhaseBuffer;  
SMUFFTHelper   *fftHelper;
```

fft size

window size

window type

```
//setup the fft
```

```
fftHelper = new SMUFFTHelper(kBufferLength, kBufferLength, WindowTypeRect);
```

```
fftMagnitudeBuffer = (float *)calloc(kBufferLength/2, sizeof(float));
```

```
fftPhaseBuffer      = (float *)calloc(kBufferLength/2, sizeof(float));
```

```
free(fftMagnitudeBuffer);
```

```
free(fftPhaseBuffer);
```

```
delete fftHelper;
```

tear down in dealloc

```
enum WindowType {  
    WindowTypeHann,  
    WindowTypeHamming,  
    WindowTypeRect,  
    WindowTypeBlackman,  
};
```

```
fftHelper->forward(0, inputAudioDataBuffer, fftMagnitudeBuffer, fftPhaseBuffer);
```

reserved


input array

magnitude out

phase out

highly optimized!! even using tricks we have not discussed

MOBILE SENSING LEARNING & CONTROL



CSE5323 & 7323

Mobile Sensing, Learning, and Control

lecture six: audio graphing, sampled data, accelerate, & FFT

Eric C. Larson, Lyle School of Engineering,
Computer Science and Engineering, Southern Methodist University